

Influence Accelerator 3.0

Project Spec — Rigorawmedia

1. What Is IA 3.0

IA 3.0 is a full-stack agentic distribution operating system. The goal is to engineer distribution — the way companies like AEOS, YaaS (YouTube as a Service), and Varun Maya's operation do it. The difference is that IA 3.0 is built from first principles up, not tactics down.

IA 2.0 was the previous version of this system. It was LLM-assisted — meaning a human ran the workflows with AI helping at each step, but no actual agents, no persistent memory, no recursive improvement. It covered pre-production, production, post-production, deployment, and attribution, but the scope was limited and a lot of it wasn't fully thought through.

IA 3.0 is the upgrade. The direction is: LLM-assisted → fully agentic → fully autonomous recursive agentic design. Every stage of the content pipeline becomes an agent. The system feeds back on itself. And the foundation is first principles, not hacks.

2. The Three-Layer Knowledge Stack

Everything in IA 3.0 is built on three layers of knowledge. This is the most important structural decision in the system because it's what prevents the whole thing from becoming a tactics pile that breaks whenever a platform changes.

Layer 1 — First Principles

This is the foundation. It's drawn from psychology, systems thinking, physics, mathematics, natural sciences, and social sciences. The mental models, heuristics, and human behavioral patterns that don't change regardless of platform or algorithm. Cognitive biases, attention mechanics, status dynamics, feedback loops — this layer is the same whether the content is a YouTube video, a LinkedIn post, or a cold email.

Layer 2 — Platform Logic

This is where first principles meet their environment. Each platform has its own constraints: algorithm behavior, content formats, retention mechanics, upload timing, what the market attention looks like at different times. The human psychology from Layer 1 is operating inside those constraints, so the application of it changes. This layer changes slowly — when platforms update their algorithm or introduce new formats, Layer 2 updates.

Layer 3 — Tactics

This is the fastest-changing layer. Hooks, trending formats, editing styles, CTA patterns, what's working right now on each platform. This layer sits on top of the other two and is constantly being updated based on what's performing. The tactics get their power from being grounded in Layers 1 and 2 — without that, they're just hacks that expire.

The platforms the system covers: LinkedIn, YouTube (long-form and short-form), Instagram (Reels and carousels), X, and newsletters. Each has its own pre-production, production, and post-production requirements.

3. The Pipeline

The full pipeline has five stages. Each stage has its own agents, tools, APIs, MCPs, and skills.

Pre-Production

This is the intelligence layer — the most important stage and the most underbuilt in most AI content systems. Pre-production is where the marketing strategy lives. It covers:

- Avatar research — deep psychological profiling of the target audience
- Research into winning content — what's already working in the market, what patterns are driving performance
- Positioning and messaging — the angle, the narrative, the specific argument being made
- Content Market Fit scoring — pre-screening content ideas before a single word is written to estimate the probability of success
- Brand spec input — every client's spec gets fed in here so the system customizes its output to that specific niche and voice

Production

Script writing, storyboarding, and actual content creation. For video this means the shoot or the AI generation. For written content it means the drafting. For short-form, HeyGen models can be used. For longer-form or avatar video at scale, the fine-tuned diffusion model is the production engine.

Post-Production

Editing, packaging, and visual polish. This is where the raw content gets turned into the final asset. Hooks get placed, thumbnails get made, captions get formatted, and the content gets packaged for its destination platform.

Deployment

Platform-specific distribution. This isn't just uploading — it's deploying at the right time, in the right format, using the best practices for that specific platform. Upload timing matters. Where market attention is weak matters. Content repurposing at scale is part of this stage too.

Metrics and Attribution

This is the feedback engine. It tracks what's actually happening after deployment — CTR, retention, shares, reach, and whatever signals matter for that platform. The attribution logic is adaptive: it updates as platforms change and as new data comes in. The whole point is that this stage feeds back into pre-production, closing the loop and making the system recursive.

The recursive loop: Research → Strategy → Create → Deploy → Measure → Learn → Update Strategy → Repeat.

4. The Content Market Fit Engine

The central innovation of IA 3.0 is the Content Market Fit (CMF) Engine. The entire goal of the Influence Accelerator system is to get content-market fit and then scale it.

The Core Problem with AI-Generated Content

AI by default produces generic output. Even with a specialized vector database or RAG setup, if the parameters given to the model are generic, the output will be generic. The problem isn't the model — it's that most systems don't extract the actual DNA behind what makes content work. They store content. They don't extract patterns.

The Three-Layer RAG System

The correct approach is a three-layer database, not a flat content dump:

Layer 1 — Raw Content Database

Scripts, videos, posts stored as-is. Most systems stop here. This alone leads to generic output because there's no abstraction — just noise.

Layer 2 — Pattern Extraction Layer

For every piece of content in the database, structured metadata gets extracted:

- Hook type
- Emotional trigger used
- Narrative structure
- Format
- Platform
- Target audience
- Outcome — views, CTR, retention, shares

This turns raw content into a structured intelligence dataset. The system isn't retrieving content — it's retrieving patterns.

Layer 3 — Judgment Layer

This is the encoded logic that tells the system which patterns matter, when to use them, for which avatar, under which conditions. This is what separates the system from a generic AI tool. The judgment layer is built from the first principles stack in Section 2 and gets refined over time as more performance data comes in.

CMF Scoring

Before any content is generated, a CMF score gets calculated. The score estimates the probability that this content will achieve market fit for this audience on this platform. The v1 formula:

- Hook Strength — 25%
- Emotional Intensity — 20%
- Novelty — 20%
- Platform Fit — 20%
- Audience Match — 15%

Output is a score out of 100 plus a recommended strategy: which hook type to use, which format, which emotional angle, which CTA pattern. The content is then generated inside those constraints, not in open-ended mode. This is what forces the AI out of generic.

Pre-Screening System

There's also a pre-screening layer that activates when a new avatar or niche is input. It auto-calculates the probability of content-market fit existing in that space — grading the market before any content is made. The system takes the full marketing strategy spec for that niche and finds the right angles, the right content strategy, the right emotional hooks for that specific engagement market.

5. The B2B Offer — AI OS for Firms

Beyond distribution and creator services, IA 3.0 powers a productized B2B implementation service. The offer is installing AI operating systems for founder-led B2B service businesses: accounting firms, consulting firms, legal practices, and recruitment agencies.

The Offer Structure

- Implementation fee: €15,000 one-time
- Monthly retainer: €2,500/month for maintenance, iteration, and support
- Target: Founder-led B2B firms with high-value workflows and no automation layer

The Technical Stack

- n8n — automation and workflow orchestration
- Lovable — rapid dashboard prototyping and client-facing UI
- MCP servers and APIs — agent execution and tool access
- Recursive agent model — extract the professional genome of the firm, bake it into agents, automate the production line
- Fine-tuned diffusion models — avatar video at scale for firms that need training content or marketing

What Gets Installed

The core idea: most B2B firms have no shortage of expertise. They have a shortage of systems. The implementation extracts what the firm already knows — its workflows, its delivery process, its client communication patterns — and encodes it into an agent system that runs without the founder's constant attention. Client delivery, outbound, content, and reporting all automated.

The 90-Day Revenue Ramp

- Days 1–7: Offer architecture, ICP definition, proposal template
- Days 8–21: 300 leads, 3 outreach sequences live, 5–8 calls booked

- Days 22–35: 2 clients closed, €30k in
 - Days 35–60: Delivery running, compounding outreach, 6–8 clients, €90–120k
 - Days 60–90: Enterprise flip, agency partnerships, productized verticals
 - Target: €250k–€1M in 90 days
-

6. Agent Architecture

Each stage of the pipeline maps to a dedicated agent or agent cluster. The build sequence is: structured prompt chains first (proto-agents), then real multi-agent orchestration, then fully autonomous recursive design.

Pre-Production Agents

- Avatar Agent — builds the psychological profile of the target audience
- Market Research Agent — pulls and analyzes winning content patterns
- Strategy Agent — generates the content strategy and messaging framework
- Content Angle Agent — produces specific content angles, hooks, and narratives

Production Agents

- Script Agent — writes scripts based on strategy output
- Story Agent — structures the narrative arc
- Format Agent — adapts content to platform-specific formats

Post-Production Agents

- Editing Agent — DaVinci MCP integration for video editing
- Packaging Agent — thumbnails, captions, visual formatting
- Hook Agent — optimizes the opening for each platform

Distribution Agents

- Platform Agent — handles deployment to YouTube, Instagram, X, LinkedIn
- Repurposing Agent — reformats content across platforms
- Scheduling Agent — deploys at optimal timing per platform

Attribution Agents

- Metrics Agent — pulls performance data from platform APIs
- Insight Agent — identifies patterns in performance data
- Strategy Update Agent — feeds insights back into pre-production parameters

Orchestration

The orchestrator coordinates all agents, maintains state, and manages the flow between stages. Tool options: n8n for automation, OpenClaw-style orchestration for agent coordination, LangGraph or CrewAI as optional layers. Different models have different strengths — Perplexity for research, Claude for writing and reasoning, other models for orchestration.

7. Obsidian Architecture

The knowledge layer lives in Obsidian. This is the brain of the system — where all the first principles, platform logic, tactics, skills, and agent definitions are stored. Every file needs to be structured as agent-readable logic, not just notes. Each file has a purpose, inputs, a process, and outputs.

Folder Structure

- /01_Principles — immutable first principles from psychology, systems, behavioral science
- /02_Platforms — platform-specific logic for YouTube, Instagram, X, LinkedIn, newsletters
- /03_Tactics — fast-changing: hooks, formats, trends, editing styles
- /04_Skills — SOPs for each agent function
- /05_Agents — agent definitions, inputs, outputs, rules
- /06_Clients — client profiles, brand specs, performance history

Skill File Format

Every skill file follows this structure:

- Name
- Input — what the agent needs to run this skill
- Process — step-by-step logic
- Output — what gets produced

Example: Hook Generation. Input: topic, audience, platform. Process: identify emotional trigger, apply curiosity gap, add pattern interrupt. Output: 5 hooks in the CMF-aligned format for that platform.

Tooling

- Obsidian — knowledge brain, principles, skills, agent logic
 - Figma — system map, flow visualization, architecture diagrams
 - n8n / agent runtime — execution layer
 - Vector DB (Pinecone or Weaviate) — pattern storage and retrieval
 - Dashboard (Lovable) — unified view across all system outputs
-

8. Cost Estimation

Two versions: the Frankenstein lean build for getting proof fast, and the enterprise build for selling to firms.

Frankenstein Version — Lean

- LLM (Claude / GPT): \$20–\$100/month
- Research (Perplexity): \$20/month
- Automation (n8n self-hosted): \$5–\$20/month
- Storage (Notion): \$10–\$20/month
- Misc APIs: free–\$20/month
- Total monthly: \$50–\$200
- Token cost per client: \$5–\$30/month
- Chargeable rate: \$500–\$2,000/month
- Build time: 10–21 days

Enterprise Version

- LLM APIs: \$100–\$500/month
- Vector DB (Pinecone / Weaviate): \$50–\$200/month
- Backend hosting: \$20–\$100/month
- Automation infra: \$20–\$50/month
- Monitoring and misc: \$30–\$100/month
- Total monthly: \$200–\$1,000
- Token cost per client: \$50–\$300/month
- Chargeable rate: \$3,000–\$15,000/month
- Build time: 6–12 weeks

The real cost in both versions is cognitive load — managing multiple agents, debugging outputs, handling ambiguity, and maintaining taste as the bottleneck. Token cost is not the bottleneck. Time is.

The correct path: build Frankenstein first, get results, understand where it breaks, then evolve into enterprise. One loop that works beats a full system that doesn't.

9. Build Order

The wrong move is jumping straight into MCPs, APIs, and complex infrastructure. The correct sequence is knowledge first, architecture second, execution third.

- Step 1 (Days 1–5): Build the knowledge layer in Obsidian — write the first principles, document platform logic for one platform, write 5–10 core skill files
- Step 2 (Days 5–7): Define all agent roles — inputs, outputs, rules for each
- Step 3 (Days 7–14): Build one complete loop on one platform. Pick YouTube or Instagram. Run the full cycle: strategy → script → produce → deploy → measure → feed back
- Step 4 (Days 14–21): Add automation via n8n and platform API connections
- Step 5 (Day 21+): Expand to more platforms and more agents

Without a clean system, agents amplify chaos. Build the foundation before building the automation.

10. The Macro Thesis

The adoption curve for AI goes: founders first, then social media creators, then enterprise and firms. Firms are the slowest to adopt. The gap between the theoretical potential of AI automation and actual real-world implementation is where the opportunity lives. That gap is the business.

The AI penetration data shows a large delta between current penetration (red) and potential penetration (blue). The question is how fast it closes. The answer is 3–5 years minimum for the fast movers, 6–7 years for broader deployment. That runway is enough to build, sell, and compound.

Even in a worst-case scenario — where model commoditization accelerates or macro shifts hit the agent economy — the distribution system is anti-fragile. It's built on immutable laws of human psychology: status, belonging, and freedom. These don't change when a new model drops. They don't change when a platform updates its algorithm. The reef survives the waves.

The evolution of the market: human only → human plus AI → AI with human oversight → fully autonomous recursive agents. IA 3.0 is being built at the inflection point between the second and third phase. That's the window.

Most builders in this space are optimizing for speed — trying to outrun model updates. IA 3.0 is optimizing for depth. First principles as the foundation, judgment encoded into the system, distribution as the moat. When the dust settles, the systems built from first principles are the ones that survive.

11. Outbound Systems

The 3 Citadel Sieging System

300 leads per week across five channels: cold email, LinkedIn, Instagram, Skool, and Loom video outreach. The system runs three sequences simultaneously across these channels for each ICP.

- Story Selling framework — narrative-led outreach that builds context before making an ask
- Nick Verge POWER Copy process — the structural framework for each message
- 88 cognitive biases library — applied at the copy level for each touchpoint
- Psychological profiling per ICP — the outreach is personalized based on a profile built before contact

Outreach Approach

The open is always a demonstrated insight, not a pitch. The prospect should feel that research has been done on their specific situation before anything is sold. The ask is small — make them curious, not convinced.

The angle: lead with a specific observation about their business. Show a quick demo or output built from that observation. Then offer a short AI intake that qualifies whether the service is even relevant for them — this removes friction and signals confidence.

Personalize deeply. Be human. Vulnerability in outreach is a differentiator, not a weakness.

12. Open Decisions

These are the questions that need answers before the next phase of the build.

- Which platform gets the first full loop? YouTube, Instagram, X, or LinkedIn?
- Which vertical gets the first client? B2B AI OS for firms, or creator distribution system?
- What is the first productized offer that can go to market in 30 days with the existing stack?
- Frankenstein version first, or sell the vision and hire delivery?
- Technical orchestration (n8n, APIs) — build in-house or outsource that layer?